Large Language Model Tutor for Global Minority Languages

Introduction

Problem Description

The recent success of Large Language Models (LLMs) has driven a great deal of interest in applications ranging from homework help to global domination.

One of the most exciting possibilities for LLMs is their application to problems of knowledge accessibility.

Companies like DuoLingo and Khan Academy are leveraging LLMs and other forms of machine learning to provide a replacement for a human tutor. However, there is a dearth of learning resources available for global minority languages, that is, languages that are not among the languages most widely spoken.

One such language is Nepali. Nepali is spoken natively by 16 million people and is used as a second language by an additional 9 million, yet it is rarely available as an option for machine translation and LLMs do not cater to its speakers.

However, before effective tools can be made for the Nepali speaking population, it must also be possible for engineers and data scientists to learn Nepali.

Project Description

In this project I will endeavor to create an AI Nepali tutor. The tutor should be able to understand English language sentences from its student and provide Nepali translations of those sentences. Additionally, it should be able to generate novel sentences in English and provide Nepali translations of those sentences.

Further, Nepali uses the Devanagari script: a form of writing unfamiliar to many English speakers, especially in the west. Thus, the student should be able to provide an unfamiliar sentence in Nepali, written in this script, and receive a translation from the tutor.

For this project I will train a series of models. The first will translate English sentences to Nepali. The second will translate written Nepali into written English. Finally, the third will generate novel sentences in English.

This project presents a unique challenge. There is very little high-quality data available for training on the Nepali language. As such, I've put significant effort into optimizing my models to be as effective as possible with with very little data.

A github repo for this project can be found here:

https://github.com/billingsmoore/nepali-tutor

Note: Your computer may have trouble rendering the Devanagari script resulting in blanks where Nepali sentences should be. Viewing this notebook as a PDF should resolve that issue.

English to Nepali Translation

First, I will create a model to translate English sentences into Nepali sentences. To create this model, I drew on the Keras tutorial provided here:

https://keras.io/examples/nlp/neural_machine_translation_with_keras_nlp/

I've adapted the model from the tutorial to translate English into Nepali, rather than Spanish, and streamlined the code for simplicity and to meet my need for computational efficiency. Additionally, I've substantially altered the model in order to more fully optimize for the much, much smaller dataset available for the Nepali language.

The first step of this process is to import the necessary libraries.

Setup

```
In [ ]: import pathlib
   import random
   import tensorflow as tf
   from tensorflow import keras
   import keras_nlp
   import matplotlib.pyplot as plt
```

Next, I will establish the necessary constants for the model.

I will use a batch size of 4 for my data. This an EXTREMELY small batch size. I'm using such a small batch size in hopes of it helping to account for the extreme smallness of my dataset. This will lead to slower convergence. However, with a larger number of epochs and a careful choice of optimization algorithms (to be discussed later) we can get away with it.

I will train the model for 100 epochs. I had initially hoped to get away with fewer, however, with such small batch sizes, convergence takes longer.

I also establish a size for the vocabulary that the model will use and the dimensions for the model to expect from the data.

An interesting addition here is AUTOTUNE. tf.data.AUTOTUNE will automate optimization for training the model. This is extremely useful both for effectively utilizing computing resources, and for avoiding too much time lost to optimization tinkering.

Importing and Exploring the Data

Now I will import the dataset that this model will be trained on. This data comes from Anki, a free, open-source flashcard program that is popular with language learners. It is available from this link:

https://www.manythings.org/anki/

I've then split the sentence pairs into two sets and set every English letter to be lowercase to avoid any confusion in the model. This is not necessary for Nepali because the Devanagari script does not use upper and lower cases.

```
In [ ]: text_file = pathlib.Path('datasets/npi-eng/npi.txt')

with open(text_file) as f:
    lines = f.read().split("\n")[:-1]

text_pairs = []

for line in lines:
    eng, nep = line.split("\t")[:2]
    eng = eng.lower()
    text_pairs.append((eng, nep))
```

The data comes in the form of numerous sentence pairs. First the sentence is given in English, then in Nepali. Each pair also has a source attribution, but that won't be necessary for the model. Below, I've printed some representative sentence pairs.

Now, we can split the sentence pairs into training, validation, and test sets. Notice that this dataset is quite small. This is one of the challenges of creating models for global minority languages. There is substantially less data to work with than if we were working with, for example, Spanish or French. As a result, I've allocated just 5% of the pairs to validation and testing respectively.

```
In []: random.shuffle(text_pairs)
    num_val_samples = int(0.05 * len(text_pairs))
    num_train_samples = len(text_pairs) - 2 * num_val_samples
    train_pairs = text_pairs[:num_train_samples]
    val_pairs = text_pairs[num_train_samples : num_train_samples + num_val_sampl
    test_pairs = text_pairs[num_train_samples + num_val_samples:]

    print(f"{len(text_pairs)} total pairs")
    print(f"{len(train_pairs)} training pairs")
    print(f"{len(val_pairs)} validation pairs")
    print(f"{len(test_pairs)} test pairs")

1574 total pairs
1418 training pairs
78 validation pairs
78 test pairs
```

Creating the Tokenizer

The tokenizer will assign each unique word in the dataset a 'token' a unique number that allows the data to be treated numerically during model training. In order to do this, a "vocabulary" must first be created. This is a complete list of the unique English and Nepali words in the dataset.

Vocabulary

```
In [ ]: def train_word_piece(text_samples, vocab_size, reserved_tokens):
    word_piece_ds = tf.data.Dataset.from_tensor_slices(text_samples)
    vocab = keras_nlp.tokenizers.compute_word_piece_vocabulary(
        word_piece_ds.batch(1000).prefetch(2),
        vocabulary_size=vocab_size,
        reserved_tokens=reserved_tokens,
```

```
return vocab
```

Tokenizing

Note that I've set aside some peculiar tokens. These correspond to whitespace,unknown characters, the beginnings and endings of sentences. I don't want the tokenizer to treat these things as words that need to be tokenized.

```
2023-08-18 10:59:12.371422: I tensorflow/compiler/xla/stream_executor/cuda/c uda_gpu_executor.cc:995] successful NUMA node read from SysFS had negative v alue (-1), but there must be at least one NUMA node, so returning NUMA node zero. See more at https://github.com/torvalds/linux/blob/v6.0/Documentation/ABI/testing/sysfs-bus-pci#L344-L355
```

2023-08-18 10:59:14.363718: I tensorflow/compiler/xla/stream_executor/cuda/c uda_gpu_executor.cc:995] successful NUMA node read from SysFS had negative v alue (-1), but there must be at least one NUMA node, so returning NUMA node zero. See more at https://github.com/torvalds/linux/blob/v6.0/Documentation/ABI/testing/sysfs-bus-pci#L344-L355

2023-08-18 10:59:14.364029: I tensorflow/compiler/xla/stream_executor/cuda/c uda_gpu_executor.cc:995] successful NUMA node read from SysFS had negative v alue (-1), but there must be at least one NUMA node, so returning NUMA node zero. See more at https://github.com/torvalds/linux/blob/v6.0/Documentation/ABI/testing/sysfs-bus-pci#L344-L355

2023-08-18 10:59:14.386559: I tensorflow/compiler/xla/stream_executor/cuda/c uda_gpu_executor.cc:995] successful NUMA node read from SysFS had negative v alue (-1), but there must be at least one NUMA node, so returning NUMA node zero. See more at https://github.com/torvalds/linux/blob/v6.0/Documentation/ABI/testing/sysfs-bus-pci#L344-L355

2023-08-18 10:59:14.387023: I tensorflow/compiler/xla/stream_executor/cuda/c uda_gpu_executor.cc:995] successful NUMA node read from SysFS had negative v alue (-1), but there must be at least one NUMA node, so returning NUMA node zero. See more at https://github.com/torvalds/linux/blob/v6.0/Documentation/ABI/testing/sysfs-bus-pci#L344-L355

2023-08-18 10:59:14.387300: I tensorflow/compiler/xla/stream_executor/cuda/c uda_gpu_executor.cc:995] successful NUMA node read from SysFS had negative v alue (-1), but there must be at least one NUMA node, so returning NUMA node zero. See more at https://github.com/torvalds/linux/blob/v6.0/Documentation/ABI/testing/sysfs-bus-pci#L344-L355

2023-08-18 10:59:14.714651: I tensorflow/compiler/xla/stream_executor/cuda/c uda_gpu_executor.cc:995] successful NUMA node read from SysFS had negative v alue (-1), but there must be at least one NUMA node, so returning NUMA node zero. See more at https://github.com/torvalds/linux/blob/v6.0/Documentation/ABI/testing/sysfs-bus-pci#L344-L355

2023-08-18 10:59:14.715049: I tensorflow/compiler/xla/stream_executor/cuda/c uda_gpu_executor.cc:995] successful NUMA node read from SysFS had negative v alue (-1), but there must be at least one NUMA node, so returning NUMA node zero. See more at https://github.com/torvalds/linux/blob/v6.0/Documentation/ABI/testing/sysfs-bus-pci#L344-L355

2023-08-18 10:59:14.715398: I tensorflow/compiler/xla/stream_executor/cuda/c uda_gpu_executor.cc:995] successful NUMA node read from SysFS had negative v alue (-1), but there must be at least one NUMA node, so returning NUMA node zero. See more at https://github.com/torvalds/linux/blob/v6.0/Documentation/ABI/testing/sysfs-bus-pci#L344-L355

2023-08-18 10:59:14.715709: I tensorflow/core/common_runtime/gpu/gpu_device. cc:1639] Created device /job:localhost/replica:0/task:0/device:GPU:0 with 10 075 MB memory: -> device: 0, name: NVIDIA GeForce RTX 4070, pci bus id: 000 0:01:00.0, compute capability: 8.9

Below we can see some example words from the dataset. Note that Nepali uses a distinct writing system that may not render correctly.

```
In [ ]: print("English Tokens: ", eng_vocab[150:155])
print("Nepali Tokens: ", nep_vocab[200:205])
```

```
English Tokens: ['much', 'old', 're', 'really', '##al']
Nepali Tokens: ['किन', 'खान', 'तिमीलाई', 'पर्छ', 'भनेर']

Finally, we can tokenize the vocabularies.

In []: eng_tokenizer = keras_nlp.tokenizers.WordPieceTokenizer(
    vocabulary=eng_vocab, lowercase=False
)

nep_tokenizer = keras_nlp.tokenizers.WordPieceTokenizer(
    vocabulary=nep_vocab, lowercase=False
)
```

Data Preprocessing

Next, I will preprocess each batch of data. This consists of re-assembling the English-Nepali sentence pairs. Each sentence must be padded with the "[PAD]" whitespace token in order to make each sequence of tokens the same length. This is because the model expects inputs of a particular shape. Once the sentence has been padded to the appropriate length, a [START] token can be appended to the beginning and an [END] token appended to the end.

Finally, this assembled dataset can be split into training and validation sets.

```
In [ ]: def eng nep preprocess batch(eng, nep):
            batch size = tf.shape(nep)[0]
            eng = eng tokenizer(eng)
            nep = nep tokenizer(nep)
            # pad eng to max sequence length
            eng start end packer = keras nlp.layers.StartEndPacker(
                sequence_length=MAX_SEQUENCE_LENGTH,
                pad value = eng tokenizer.token to id("[PAD]"),
            eng = eng start end packer(eng)
            # add special tokens [start] and [end] and pad nep
            nep start end packer = keras nlp.layers.StartEndPacker(
                sequence length = MAX SEQUENCE LENGTH + 1,
                start value = nep tokenizer.token to id("[START]"),
                end_value = nep_tokenizer.token_to_id("[END]"),
                pad value = nep tokenizer.token to id("[PAD]")
            nep = nep start end packer(nep)
            return (
                "encoder inputs": eng,
                "decoder inputs": nep[:, :-1]
                },
```

Creating the Model

Now it's time to build the model itself. This model is an Autoencoder, which consists of an encoder and a decoder.

The encoder input layer takes in a set of tokenized inputs. These inputs are then passed to a layer that accounts for the number assigned to the token as well as the position of that token in the sentence. The next layer is a typical dense Encoder layer.

The decoder takes in a set of tokenized inputs from the Nepali dataset and passes them to a layer that will account for the token number and position of the token in those sentences. This is then passed to a typical dense Decoder layer.

Both the Encoder and Decoder layers are helpfully provided out-of-the-box by Keras.

```
In [ ]: encoder inputs = keras.Input(shape=(None,), dtype="int64", name="encoder input
        x = keras nlp.layers.TokenAndPositionEmbedding(
            vocabulary size=VOCAB SIZE,
            sequence length = MAX SEQUENCE LENGTH,
            embedding dim=EMBED DIM,
            mask zero=True,
        )(encoder inputs)
        encoder outputs = keras nlp.layers.TransformerEncoder(
            intermediate dim = INTERMEDIATE DIM, num heads=NUM HEADS
        )(inputs=x)
        encoder = keras.Model(encoder inputs, encoder outputs)
        decoder inputs = keras.Input(shape=(None,), dtype="int64", name="decoder inp
        encoded seq inputs = keras.Input(shape=(None, EMBED DIM), name="decoder stat
        x = keras nlp.layers.TokenAndPositionEmbedding(
            vocabulary size=VOCAB SIZE,
            sequence length=MAX SEQUENCE LENGTH,
            embedding dim=EMBED DIM,
            mask zero=True,
        )(decoder inputs)
```

```
x = keras_nlp.layers.TransformerDecoder(
    intermediate_dim=INTERMEDIATE_DIM, num_heads=NUM_HEADS
)(decoder_sequence=x, encoder_sequence=encoded_seq_inputs)
x = keras.layers.Dropout(0.5)(x)
decoder_outputs = keras.layers.Dense(VOCAB_SIZE, activation="softmax")(x)
decoder = keras.Model(
    [
         decoder_inputs,
         encoded_seq_inputs
],
    decoder_outputs = decoder([decoder_inputs, encoder_outputs])
eng_nep_translator = keras.Model(
    [encoder_inputs, decoder_inputs],
    decoder_outputs,
    name="eng_nep_translator",
)
```

Model Summary

```
In [ ]: eng_nep_translator.summary()
```

9 of 41

Model: "eng_nep_translator"

Layer (type) d to	Output Shape	Param #	Connecte
encoder_inputs (InputLayer)	[(None, None)]	0	[]
Layer (type) d to	Output Shape	Param #	Connecte
<pre>====================================</pre>	[(None, None)]	0	[]
<pre>token_and_position_embeddi r_inputs[0][0]'] ng (TokenAndPositionEmbedd ing)</pre>	(None, None, 256)	3850240	['encode
<pre>decoder_inputs (InputLayer)</pre>	[(None, None)]	0	[]
<pre>transformer_encoder (Trans and_position_embedding formerEncoder)</pre>	(None, None, 256)	1315072	['token_ [0][0]']
<pre>model_1 (Functional) r_inputs[0][0]',</pre>	(None, None, 15000)	9283992	['decode
ormer_encoder[0][0]']			'transf
Total params: 14449304 (55.12 MB) Trainable params: 14449304 (55.12 MB) Non-trainable params: 0 (0.00 Byte)			

Compilation

Now, I've compiled the model.

Of note here is the choice of optimization algorith. I have used RMSProp. RMSProp is similar to Adagrad, which we studied in class, and as a result it converges much more quickly than, say, SGD. However, it is less susceptible to vanishing gradients. This is perfect for our small dataset with small batch sizes.

The loss function is Sparse Categorical Crossentropy. Not every word appears in every

sentence so the data for most natural language related tasks is necessarily sparse.

```
In [ ]: eng_nep_translator.compile(
         "rmsprop",
         loss="sparse_categorical_crossentropy",
         metrics=["accuracy"]
)
```

Fitting the Model

```
Epoch 1/100
23/23 [============== ] - 15s 301ms/step - loss: 5.7427 - acc
uracy: 0.1268 - val loss: 4.6218 - val accuracy: 0.2142
23/23 [============== ] - Os 21ms/step - loss: 4.3512 - accur
acy: 0.2179 - val loss: 4.2112 - val accuracy: 0.2256
23/23 [============== ] - 1s 24ms/step - loss: 3.9231 - accur
acy: 0.2553 - val loss: 3.8240 - val accuracy: 0.2549
Epoch 4/100
23/23 [============== ] - 1s 22ms/step - loss: 3.6029 - accur
acy: 0.2820 - val loss: 3.5899 - val accuracy: 0.2769
Epoch 5/100
23/23 [============== ] - 1s 22ms/step - loss: 3.3899 - accur
acy: 0.2981 - val loss: 3.5028 - val accuracy: 0.2834
Epoch 6/100
acy: 0.3088 - val loss: 3.4117 - val accuracy: 0.2818
Epoch 7/100
23/23 [============== ] - Os 20ms/step - loss: 3.1218 - accur
acy: 0.3177 - val loss: 3.3454 - val accuracy: 0.2956
Epoch 8/100
23/23 [============ ] - 0s 21ms/step - loss: 3.0220 - accur
acy: 0.3251 - val loss: 3.2250 - val accuracy: 0.3119
Epoch 9/100
23/23 [============== ] - 1s 24ms/step - loss: 2.9113 - accur
acy: 0.3407 - val loss: 3.2453 - val accuracy: 0.3111
Epoch 10/100
23/23 [============== ] - 0s 20ms/step - loss: 2.8396 - accur
acy: 0.3457 - val_loss: 3.2248 - val_accuracy: 0.3200
Epoch 11/100
acy: 0.3575 - val loss: 3.0920 - val accuracy: 0.3217
Epoch 12/100
acy: 0.3706 - val loss: 3.0824 - val accuracy: 0.3135
Epoch 13/100
23/23 [=============== ] - 0s 20ms/step - loss: 2.6100 - accur
acy: 0.3795 - val loss: 3.0974 - val accuracy: 0.3331
Epoch 14/100
23/23 [============ ] - 0s 21ms/step - loss: 2.5298 - accur
acy: 0.3945 - val_loss: 3.0074 - val_accuracy: 0.3388
Epoch 15/100
acy: 0.4047 - val loss: 2.9844 - val accuracy: 0.3436
Epoch 16/100
23/23 [============ ] - 0s 20ms/step - loss: 2.3783 - accur
acy: 0.4193 - val loss: 3.0206 - val accuracy: 0.3445
Epoch 17/100
23/23 [============== ] - 0s 20ms/step - loss: 2.3049 - accur
acy: 0.4322 - val loss: 2.9466 - val accuracy: 0.3575
Epoch 18/100
23/23 [============== ] - 1s 23ms/step - loss: 2.2295 - accur
acy: 0.4487 - val_loss: 2.9362 - val_accuracy: 0.3648
```

```
acy: 0.4660 - val loss: 2.8849 - val accuracy: 0.3599
Epoch 20/100
acy: 0.4803 - val loss: 2.9393 - val accuracy: 0.3640
Epoch 21/100
acy: 0.5017 - val loss: 2.9228 - val accuracy: 0.3770
Epoch 22/100
23/23 [============== ] - 0s 20ms/step - loss: 1.8585 - accur
acy: 0.5265 - val loss: 2.8914 - val accuracy: 0.3884
Epoch 23/100
23/23 [============== ] - 0s 21ms/step - loss: 1.7961 - accur
acy: 0.5374 - val loss: 2.8964 - val accuracy: 0.3958
Epoch 24/100
acy: 0.5622 - val loss: 2.8951 - val accuracy: 0.4055
Epoch 25/100
23/23 [============== ] - 0s 22ms/step - loss: 1.6087 - accur
acy: 0.5791 - val loss: 2.9280 - val_accuracy: 0.4047
Epoch 26/100
acy: 0.6144 - val loss: 2.9629 - val accuracy: 0.3974
Epoch 27/100
23/23 [============== ] - 1s 23ms/step - loss: 1.3917 - accur
acy: 0.6302 - val loss: 2.9667 - val accuracy: 0.4169
Epoch 28/100
acy: 0.6569 - val loss: 3.0268 - val accuracy: 0.4153
Epoch 29/100
23/23 [============= ] - Os 21ms/step - loss: 1.1906 - accur
acy: 0.6793 - val loss: 3.0415 - val accuracy: 0.4259
Epoch 30/100
23/23 [============== ] - 1s 24ms/step - loss: 1.0887 - accur
acy: 0.7080 - val loss: 3.0412 - val accuracy: 0.4169
Epoch 31/100
23/23 [============== ] - 0s 21ms/step - loss: 1.0141 - accur
acy: 0.7255 - val loss: 3.1149 - val accuracy: 0.4235
Epoch 32/100
acy: 0.7504 - val loss: 3.1392 - val accuracy: 0.4137
Epoch 33/100
acy: 0.7783 - val_loss: 3.1473 - val accuracy: 0.4340
Epoch 34/100
23/23 [============== ] - 0s 21ms/step - loss: 0.7574 - accur
acy: 0.7898 - val loss: 3.2356 - val accuracy: 0.4267
Epoch 35/100
acy: 0.8183 - val_loss: 3.2747 - val_accuracy: 0.4316
Epoch 36/100
23/23 [============== ] - 0s 20ms/step - loss: 0.6310 - accur
acy: 0.8265 - val loss: 3.3351 - val accuracy: 0.4373
Epoch 37/100
acy: 0.8472 - val loss: 3.3859 - val accuracy: 0.4324
Epoch 38/100
```

```
23/23 [============== ] - 1s 23ms/step - loss: 0.4769 - accur
acy: 0.8699 - val loss: 3.4272 - val accuracy: 0.4283
Epoch 39/100
23/23 [============== ] - 0s 21ms/step - loss: 0.4445 - accur
acy: 0.8798 - val_loss: 3.4757 - val_accuracy: 0.4454
Epoch 40/100
23/23 [============= ] - Os 19ms/step - loss: 0.3955 - accur
acy: 0.8925 - val loss: 3.5336 - val accuracy: 0.4292
Epoch 41/100
23/23 [============== ] - 1s 23ms/step - loss: 0.3527 - accur
acy: 0.9047 - val loss: 3.5205 - val accuracy: 0.4389
Epoch 42/100
23/23 [============== ] - 1s 22ms/step - loss: 0.3362 - accur
acy: 0.9081 - val loss: 3.6622 - val accuracy: 0.4332
Epoch 43/100
23/23 [============= ] - 1s 22ms/step - loss: 0.3422 - accur
acy: 0.9060 - val loss: 3.6460 - val accuracy: 0.4430
Epoch 44/100
acy: 0.9253 - val_loss: 3.7569 - val_accuracy: 0.4300
Epoch 45/100
23/23 [============== ] - 0s 20ms/step - loss: 0.2509 - accur
acy: 0.9306 - val loss: 3.8573 - val accuracy: 0.4251
Epoch 46/100
23/23 [============ ] - 1s 22ms/step - loss: 0.2338 - accur
acy: 0.9372 - val loss: 3.8833 - val accuracy: 0.4381
Epoch 47/100
23/23 [============= ] - 1s 24ms/step - loss: 0.2231 - accur
acy: 0.9386 - val loss: 3.8998 - val accuracy: 0.4397
Epoch 48/100
23/23 [============== ] - 0s 21ms/step - loss: 0.2155 - accur
acy: 0.9408 - val loss: 3.9011 - val accuracy: 0.4438
Epoch 49/100
23/23 [============= ] - 0s 20ms/step - loss: 0.1991 - accur
acy: 0.9461 - val loss: 3.9741 - val accuracy: 0.4422
Epoch 50/100
23/23 [============== ] - 1s 23ms/step - loss: 0.1889 - accur
acy: 0.9498 - val loss: 4.0333 - val accuracy: 0.4316
Epoch 51/100
23/23 [============== ] - Os 21ms/step - loss: 0.1833 - accur
acy: 0.9500 - val loss: 4.0515 - val accuracy: 0.4349
Epoch 52/100
23/23 [============= ] - Os 20ms/step - loss: 0.1658 - accur
acy: 0.9549 - val loss: 4.0739 - val accuracy: 0.4340
Epoch 53/100
acy: 0.9573 - val loss: 4.1424 - val accuracy: 0.4332
Epoch 54/100
acy: 0.9560 - val loss: 4.1376 - val accuracy: 0.4454
Epoch 55/100
23/23 [============ ] - 0s 21ms/step - loss: 0.1418 - accur
acy: 0.9612 - val loss: 4.1774 - val accuracy: 0.4389
Epoch 56/100
23/23 [============== ] - 1s 23ms/step - loss: 0.1401 - accur
acy: 0.9647 - val loss: 4.1605 - val accuracy: 0.4446
```

```
Epoch 57/100
23/23 [============== ] - 0s 20ms/step - loss: 0.1364 - accur
acy: 0.9643 - val loss: 4.2501 - val accuracy: 0.4275
Epoch 58/100
23/23 [============= ] - Os 21ms/step - loss: 0.1316 - accur
acy: 0.9652 - val loss: 4.2439 - val accuracy: 0.4536
Epoch 59/100
23/23 [============== ] - 1s 23ms/step - loss: 0.1282 - accur
acy: 0.9654 - val loss: 4.2698 - val accuracy: 0.4463
Epoch 60/100
acy: 0.9675 - val loss: 4.3258 - val accuracy: 0.4414
Epoch 61/100
23/23 [============== ] - 0s 20ms/step - loss: 0.1196 - accur
acy: 0.9679 - val loss: 4.3740 - val accuracy: 0.4454
Epoch 62/100
23/23 [============= ] - 1s 22ms/step - loss: 0.1154 - accur
acy: 0.9695 - val loss: 4.3361 - val accuracy: 0.4389
Epoch 63/100
23/23 [============= ] - Os 21ms/step - loss: 0.1115 - accur
acy: 0.9700 - val loss: 4.4678 - val accuracy: 0.4365
Epoch 64/100
23/23 [============ ] - 0s 20ms/step - loss: 0.1155 - accur
acy: 0.9684 - val loss: 4.4577 - val accuracy: 0.4324
Epoch 65/100
23/23 [============== ] - 1s 23ms/step - loss: 0.1036 - accur
acy: 0.9726 - val loss: 4.4356 - val accuracy: 0.4349
Epoch 66/100
23/23 [============= ] - Os 21ms/step - loss: 0.1136 - accur
acy: 0.9692 - val_loss: 4.4343 - val_accuracy: 0.4340
Epoch 67/100
23/23 [============= ] - Os 21ms/step - loss: 0.0921 - accur
acy: 0.9760 - val loss: 4.4924 - val accuracy: 0.4243
Epoch 68/100
acy: 0.9703 - val loss: 4.5033 - val accuracy: 0.4438
acy: 0.9755 - val loss: 4.5548 - val accuracy: 0.4389
Epoch 70/100
23/23 [============== ] - 0s 21ms/step - loss: 0.0907 - accur
acy: 0.9762 - val_loss: 4.4792 - val_accuracy: 0.4332
Epoch 71/100
acy: 0.9789 - val loss: 4.5505 - val accuracy: 0.4414
Epoch 72/100
23/23 [============ ] - 0s 20ms/step - loss: 0.0963 - accur
acy: 0.9757 - val loss: 4.5601 - val accuracy: 0.4357
Epoch 73/100
23/23 [============== ] - 0s 21ms/step - loss: 0.0932 - accur
acy: 0.9745 - val loss: 4.6776 - val accuracy: 0.4316
Epoch 74/100
23/23 [============== ] - 0s 20ms/step - loss: 0.0986 - accur
acy: 0.9733 - val loss: 4.5429 - val accuracy: 0.4324
```

```
acy: 0.9772 - val_loss: 4.5644 - val accuracy: 0.4389
Epoch 76/100
23/23 [============= ] - 1s 22ms/step - loss: 0.0760 - accur
acy: 0.9817 - val loss: 4.5976 - val accuracy: 0.4406
Epoch 77/100
acy: 0.9795 - val loss: 4.6826 - val accuracy: 0.4397
Epoch 78/100
23/23 [============== ] - 0s 21ms/step - loss: 0.0825 - accur
acy: 0.9785 - val loss: 4.6726 - val accuracy: 0.4544
Epoch 79/100
23/23 [============== ] - 1s 25ms/step - loss: 0.0804 - accur
acy: 0.9801 - val loss: 4.7449 - val accuracy: 0.4454
acy: 0.9796 - val loss: 4.8342 - val accuracy: 0.4389
Epoch 81/100
23/23 [============== ] - 0s 21ms/step - loss: 0.0810 - accur
acy: 0.9773 - val loss: 4.7415 - val accuracy: 0.4332
Epoch 82/100
acy: 0.9818 - val loss: 4.7434 - val accuracy: 0.4308
Epoch 83/100
23/23 [============== ] - 0s 21ms/step - loss: 0.0698 - accur
acy: 0.9819 - val loss: 4.8233 - val accuracy: 0.4430
Epoch 84/100
23/23 [============= ] - Os 21ms/step - loss: 0.0779 - accur
acy: 0.9788 - val loss: 4.7662 - val accuracy: 0.4349
Epoch 85/100
23/23 [============= ] - 1s 24ms/step - loss: 0.0780 - accur
acy: 0.9792 - val loss: 4.8281 - val accuracy: 0.4349
23/23 [============= ] - Os 20ms/step - loss: 0.0723 - accur
acy: 0.9808 - val loss: 4.7724 - val accuracy: 0.4471
Epoch 87/100
23/23 [============== ] - 0s 21ms/step - loss: 0.0709 - accur
acy: 0.9810 - val loss: 4.8149 - val accuracy: 0.4430
Epoch 88/100
23/23 [============== ] - 1s 23ms/step - loss: 0.0728 - accur
acy: 0.9800 - val loss: 4.7687 - val accuracy: 0.4487
Epoch 89/100
acy: 0.9834 - val loss: 4.8348 - val accuracy: 0.4267
Epoch 90/100
23/23 [============== ] - 0s 21ms/step - loss: 0.0703 - accur
acy: 0.9818 - val loss: 4.8870 - val accuracy: 0.4389
Epoch 91/100
acy: 0.9815 - val_loss: 4.8855 - val_accuracy: 0.4365
Epoch 92/100
23/23 [============== ] - 0s 21ms/step - loss: 0.0644 - accur
acy: 0.9830 - val loss: 4.8516 - val accuracy: 0.4430
Epoch 93/100
acy: 0.9820 - val loss: 4.8644 - val accuracy: 0.4422
Epoch 94/100
```

```
acy: 0.9816 - val loss: 4.8650 - val accuracy: 0.4397
Epoch 95/100
acy: 0.9826 - val_loss: 4.8566 - val_accuracy: 0.4300
Epoch 96/100
23/23 [============== ] - 0s 21ms/step - loss: 0.0629 - accur
acy: 0.9832 - val_loss: 4.8677 - val_accuracy: 0.4422
Epoch 97/100
23/23 [============== ] - 0s 21ms/step - loss: 0.0682 - accur
acy: 0.9816 - val loss: 4.8726 - val accuracy: 0.4463
Epoch 98/100
acy: 0.9837 - val loss: 4.7683 - val accuracy: 0.4430
Epoch 99/100
23/23 [============== ] - 1s 23ms/step - loss: 0.0597 - accur
acy: 0.9852 - val loss: 4.7806 - val accuracy: 0.4389
Epoch 100/100
acy: 0.9854 - val_loss: 4.9264 - val_accuracy: 0.4397
```

To avoid training and retraining the model, I'll now save this model with these results.

Visualizing the Training Results

Below, we can see how the loss and accuracy evolved over the course of training. Here we can really see how difficult it is to make effective generative tools from small datasets. Even as the model's accuracy improves substantially on the training set, the accuracy on the validation data remains unacceptably low.

The loss on the validation data also never decreases, instead getting worse as time goes on.

```
In []: acc = eng_nep_history.history['accuracy']
    val_acc = eng_nep_history.history['val_accuracy']

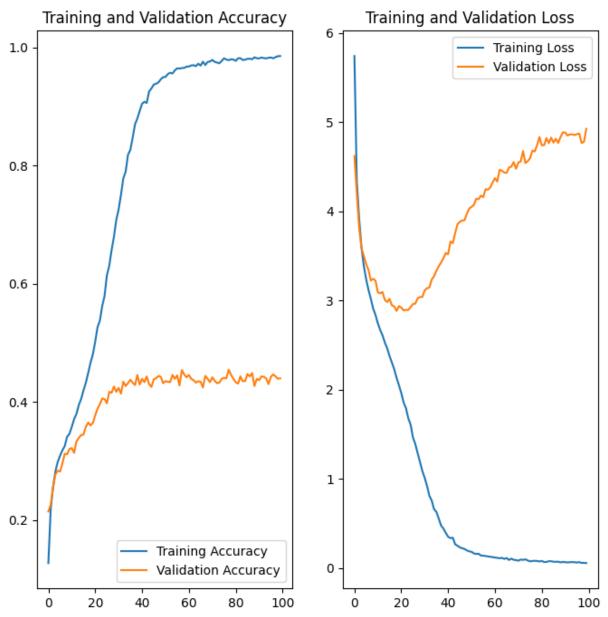
loss = eng_nep_history.history['loss']
    val_loss = eng_nep_history.history['val_loss']

epochs_range = range(100)

plt.figure(figsize=(8, 8))
    plt.subplot(1, 2, 1)
    plt.plot(epochs_range, acc, label='Training Accuracy')
    plt.plot(epochs_range, val_acc, label='Validation Accuracy')
    plt.legend(loc='lower right')
```

```
plt.title('Training and Validation Accuracy')

plt.subplot(1, 2, 2)
plt.plot(epochs_range, loss, label='Training Loss')
plt.plot(epochs_range, val_loss, label='Validation Loss')
plt.legend(loc='upper right')
plt.title('Training and Validation Loss')
plt.show()
```



Decoding Translated Sentences

Even if the translations are perfect, the outputs of our model are not meaningful sentences. The model only outputs numerical tokens. In order to turn these into something that a human can read they need to be decoded.

Below is a function to decode these translated sentences. This function takes in an English sentence, runs it through our translator model then works its way through the output of

the model, converting the output into words in Nepali using our tokenizers.

Part of decoding the sequence is sampling the probabilities of tokens that should follow the existing translation. The sampler is the algorithm that is used to select that next work or token. Here I've used the Greedy sampler which simply finds the highest likelihood next word and adds it to the translated sentence. It is computationally inexpensive and because the outputs are pretty short we don't need to worry about the Greedy sampler outputting long, repetitive sentences that don't make much sense, which can be an issue with the algorithm.

```
In [ ]: def eng nep translate(input sentences):
            batch size = tf.shape(input sentences)[0]
            encoder_input_tokens = eng_tokenizer(input_sentences).to tensor(
                shape=(None, MAX SEQUENCE LENGTH)
            def next(prompt, cache, index):
                logits = eng nep translator([encoder input tokens, prompt])[:, index
                hidden states = None
                return logits, hidden states, cache
            length = 40
            start = tf.fill((batch size, 1), nep tokenizer.token to id("[START]"))
            pad = tf.fill((batch size, length - 1), nep tokenizer.token to id("[PAD]
            prompt = tf.concat((start, pad), axis=-1)
            generated tokens = keras nlp.samplers.GreedySampler()(
                next,
                prompt,
                end token id=nep tokenizer.token to id("[END]"),
            generated sentences = nep tokenizer.detokenize(generated tokens)
            return generated sentences
```

Example Translations

Now, let's look at some example translations from the model.

```
In []: test_eng_texts = [pair[0] for pair in test_pairs]
for i in range(5):
    input_sentence = random.choice(test_eng_texts)
    translated = eng_nep_translate(tf.constant([input_sentence]))
    translated = translated.numpy()[0].decode("utf-8")
    translated = (
        translated.replace("[PAD]", "")
        .replace("[START]", "")
        .replace("[END]", "")
        .strip()
)
```

```
print(f"** Example {i} **")
     print(input sentence)
     print(translated)
     print()
2023-08-18 10:59:35.830764: I tensorflow/compiler/xla/stream executor/cuda/c
uda blas.cc:606] TensorFloat-32 will be used for the matrix multiplication.
This will only be logged once.
2023-08-18 10:59:36.980078: I tensorflow/compiler/xla/service/service.cc:16
8] XLA service 0x1051b930 initialized for platform CUDA (this does not guara
ntee that XLA will be used). Devices:
2023-08-18 10:59:36.980152: I tensorflow/compiler/xla/service/service.cc:17
     StreamExecutor device (0): NVIDIA GeForce RTX 4070, Compute Capability
8.9
2023-08-18 10:59:38.690555: I tensorflow/compiler/xla/stream executor/cuda/c
uda dnn.cc:432] Loaded cuDNN version 8904
2023-08-18 10:59:40.544634: I tensorflow/tsl/platform/default/subprocess.cc:
304] Start cannot spawn child process: No such file or directory
2023-08-18 10:59:41.517654: I ./tensorflow/compiler/jit/device compiler.h:18
6] Compiled cluster using XLA! This line is logged at most once for the lif
etime of the process.
** Example 0 **
i didn't want to go, but i did.
ो ठ चाहन्छ्रौं एक अहिलेला ,ब छिन्ौं ड हुँ हुनुहुन्छान थाहाको
** Example 1 **
i know that tom is in bed.
मैले दहरू ौलाई काम थाहाको
** Example 2 **
you won't see me again.
ह मैले ी ा अहिले टमलाई छिन्ो ्ईस थाहा , मैलेलेको
** Example 3 **
when was the last time you cried?
गर्न0 हामी ढसदैनिएको किस हुँै ?
** Example 4 **
```

Nepali to English Translation

i'm pretty sure tom's happy in boston. मैले ु चाहन्छुरही ढससँगससँगिएको कि थ एक ठदैन तपाई दको

Now that the English to Nepali Translator is finished, it's time to reverse the process.

There's no need to repeat the setup stages nor the importing and preprocessing of data, since it can all be reused from the first translator. We can simply refactor our preprocess_batch function to reverse which language is the input and which language is the output.

```
In [ ]: def nep_eng_preprocess_batch(eng, nep):
```

```
eng = eng_tokenizer(eng)
    nep = nep tokenizer(nep)
    # pad eng to max sequence length
    eng_start_end_packer = keras_nlp.layers.StartEndPacker(
        sequence length=MAX SEQUENCE LENGTH+1,
        pad value = eng tokenizer.token to id("[PAD]"),
    eng = eng start end packer(eng)
    # add special tokens [start] and [end] and pad nep
    nep start end packer = keras nlp.layers.StartEndPacker(
        sequence length = MAX SEQUENCE LENGTH,
        start value = nep tokenizer.token to id("[START]"),
        end value = nep tokenizer.token to id("[END]"),
        pad value = nep tokenizer.token to id("[PAD]")
    )
    nep = nep start end packer(nep)
    return (
        "encoder inputs": nep,
        "decoder_inputs": eng[:, :-1]
        },
        eng[:, 1:],
    )
def make_dataset(pairs):
    eng texts, nep texts = zip(*pairs)
    eng_texts = list(eng_texts)
    nep texts = list(nep texts)
    dataset = tf.data.Dataset.from tensor slices((nep texts, eng texts))
    dataset=dataset.batch(BATCH SIZE)
    dataset = dataset.map(nep_eng_preprocess_batch, num_parallel_calls=AUTOT
    return dataset.shuffle(2048).prefetch(16).cache()
nep eng train ds = make dataset(train pairs)
nep eng val ds = make dataset(val pairs)
```

Creating the Model

We can reuse the same model architecture from above.

Model Summary

```
In [ ]: nep eng translator.summary()
      Model: "nep eng translator"
       Layer (type)
                                Output Shape
                                                          Param #
                                                                   Connecte
      _____
       encoder inputs (InputLayer [(None, None)]
                                                          0
                                                                   []
       Layer (type)
                                Output Shape
                                                          Param #
                                                                   Connecte
      d to
       encoder inputs (InputLayer [(None, None)]
                                                                   []
       token and position embeddi
                                (None, None, 256)
                                                          3850240
                                                                   ['encode
      r inputs[0][0]']
       ng (TokenAndPositionEmbedd
       ing)
       decoder inputs (InputLayer [(None, None)]
                                                                   []
                                                          0
       transformer_encoder (Trans (None, None, 256)
                                                          1315072
                                                                   ['token_
      and position embedding
       formerEncoder)
                                                                   [0][0]']
                               (None, None, 15000)
       model 1 (Functional)
                                                          9283992
                                                                   ['decode
      r inputs[0][0]',
                                                                    'transf
      ormer_encoder[0][0]']
      ______
      _____
      Total params: 14449304 (55.12 MB)
      Trainable params: 14449304 (55.12 MB)
      Non-trainable params: 0 (0.00 Byte)
```

Compilation

Here I've used the same loss and optimization algorithms as the other translator.

22 of 41

```
In [ ]: nep_eng_translator.compile(
          "rmsprop",
          loss="sparse_categorical_crossentropy",
          metrics=["accuracy"]
)
```

Fitting the Model

23 of 41

```
Epoch 1/100
uracy: 0.7730 - val loss: 0.6445 - val accuracy: 0.7848
23/23 [============== ] - 1s 24ms/step - loss: 0.5696 - accur
acy: 0.8187 - val loss: 0.5252 - val accuracy: 0.8135
acy: 0.8355 - val loss: 0.4863 - val accuracy: 0.8463
Epoch 4/100
acy: 0.8464 - val loss: 0.4911 - val accuracy: 0.8135
Epoch 5/100
23/23 [============== ] - 1s 23ms/step - loss: 0.4292 - accur
acy: 0.8491 - val loss: 0.4543 - val accuracy: 0.8484
Epoch 6/100
23/23 [============== ] - Os 20ms/step - loss: 0.4058 - accur
acy: 0.8572 - val loss: 0.4350 - val accuracy: 0.8525
Epoch 7/100
23/23 [============= ] - Os 21ms/step - loss: 0.3929 - accur
acy: 0.8598 - val loss: 0.3791 - val accuracy: 0.8607
Epoch 8/100
23/23 [============ ] - 1s 26ms/step - loss: 0.3552 - accur
acy: 0.8600 - val loss: 0.3576 - val accuracy: 0.8689
Epoch 9/100
23/23 [============== ] - 0s 20ms/step - loss: 0.3352 - accur
acy: 0.8639 - val loss: 0.3597 - val accuracy: 0.8709
Epoch 10/100
23/23 [============== ] - 1s 24ms/step - loss: 0.3233 - accur
acy: 0.8699 - val_loss: 0.3577 - val_accuracy: 0.8668
Epoch 11/100
23/23 [============== ] - 0s 21ms/step - loss: 0.3192 - accur
acy: 0.8716 - val loss: 0.3513 - val accuracy: 0.8586
Epoch 12/100
acy: 0.8706 - val loss: 0.3578 - val accuracy: 0.8709
Epoch 13/100
23/23 [============== ] - 0s 20ms/step - loss: 0.2982 - accur
acy: 0.8702 - val loss: 0.3527 - val accuracy: 0.8709
Epoch 14/100
23/23 [============== ] - 1s 23ms/step - loss: 0.2976 - accur
acy: 0.8712 - val_loss: 0.3507 - val_accuracy: 0.8668
Epoch 15/100
acy: 0.8716 - val loss: 0.3481 - val accuracy: 0.8566
Epoch 16/100
23/23 [============ ] - 1s 22ms/step - loss: 0.2849 - accur
acy: 0.8758 - val loss: 0.3461 - val accuracy: 0.8566
Epoch 17/100
acy: 0.8780 - val loss: 0.3553 - val accuracy: 0.8463
Epoch 18/100
23/23 [============== ] - 0s 21ms/step - loss: 0.2785 - accur
acy: 0.8763 - val loss: 0.3430 - val accuracy: 0.8607
```

```
acy: 0.8756 - val_loss: 0.3332 - val accuracy: 0.8730
Epoch 20/100
23/23 [============= ] - 1s 23ms/step - loss: 0.2760 - accur
acy: 0.8749 - val loss: 0.3532 - val accuracy: 0.8607
Epoch 21/100
acy: 0.8786 - val loss: 0.3422 - val accuracy: 0.8730
Epoch 22/100
23/23 [============== ] - 1s 24ms/step - loss: 0.2705 - accur
acy: 0.8773 - val loss: 0.3471 - val accuracy: 0.8689
Epoch 23/100
23/23 [============== ] - 0s 21ms/step - loss: 0.2681 - accur
acy: 0.8790 - val loss: 0.3649 - val accuracy: 0.8566
Epoch 24/100
23/23 [============= ] - 0s 21ms/step - loss: 0.2703 - accur
acy: 0.8772 - val loss: 0.3464 - val accuracy: 0.8668
Epoch 25/100
23/23 [============== ] - 1s 23ms/step - loss: 0.2663 - accur
acy: 0.8780 - val loss: 0.3523 - val accuracy: 0.8627
Epoch 26/100
acy: 0.8775 - val loss: 0.3432 - val accuracy: 0.8730
Epoch 27/100
23/23 [============== ] - 1s 24ms/step - loss: 0.2639 - accur
acy: 0.8801 - val loss: 0.3502 - val accuracy: 0.8689
Epoch 28/100
23/23 [============= ] - Os 21ms/step - loss: 0.2625 - accur
acy: 0.8791 - val loss: 0.3513 - val accuracy: 0.8668
Epoch 29/100
23/23 [============== ] - 1s 24ms/step - loss: 0.2583 - accur
acy: 0.8815 - val loss: 0.3720 - val accuracy: 0.8648
Epoch 30/100
23/23 [============== ] - Os 20ms/step - loss: 0.2596 - accur
acy: 0.8781 - val loss: 0.3625 - val accuracy: 0.8668
Epoch 31/100
23/23 [============== ] - Os 21ms/step - loss: 0.2626 - accur
acy: 0.8778 - val loss: 0.3449 - val accuracy: 0.8607
Epoch 32/100
acy: 0.8827 - val loss: 0.3527 - val accuracy: 0.8648
Epoch 33/100
acy: 0.8825 - val loss: 0.3646 - val accuracy: 0.8525
Epoch 34/100
23/23 [============== ] - 0s 21ms/step - loss: 0.2576 - accur
acy: 0.8795 - val loss: 0.3614 - val accuracy: 0.8668
Epoch 35/100
acy: 0.8829 - val_loss: 0.3609 - val_accuracy: 0.8648
Epoch 36/100
23/23 [============== ] - 1s 23ms/step - loss: 0.2550 - accur
acy: 0.8852 - val loss: 0.3523 - val accuracy: 0.8607
Epoch 37/100
acy: 0.8832 - val loss: 0.3474 - val accuracy: 0.8627
Epoch 38/100
```

```
acy: 0.8842 - val loss: 0.3473 - val accuracy: 0.8668
Epoch 39/100
23/23 [============== ] - 1s 24ms/step - loss: 0.2467 - accur
acy: 0.8834 - val_loss: 0.3717 - val_accuracy: 0.8586
Epoch 40/100
acy: 0.8831 - val_loss: 0.3653 - val_accuracy: 0.8627
Epoch 41/100
23/23 [============== ] - 0s 20ms/step - loss: 0.2448 - accur
acy: 0.8855 - val loss: 0.3647 - val accuracy: 0.8627
Epoch 42/100
23/23 [============== ] - 1s 23ms/step - loss: 0.2445 - accur
acy: 0.8860 - val loss: 0.3688 - val accuracy: 0.8504
Epoch 43/100
23/23 [============== ] - 0s 20ms/step - loss: 0.2483 - accur
acy: 0.8850 - val loss: 0.3493 - val accuracy: 0.8648
Epoch 44/100
acy: 0.8832 - val loss: 0.3494 - val accuracy: 0.8545
Epoch 45/100
23/23 [============== ] - 1s 23ms/step - loss: 0.2465 - accur
acy: 0.8837 - val loss: 0.3739 - val accuracy: 0.8648
Epoch 46/100
23/23 [============= ] - 0s 21ms/step - loss: 0.2514 - accur
acy: 0.8833 - val loss: 0.3594 - val accuracy: 0.8689
Epoch 47/100
23/23 [============ ] - 0s 22ms/step - loss: 0.2447 - accur
acy: 0.8836 - val loss: 0.3539 - val accuracy: 0.8648
Epoch 48/100
23/23 [============== ] - 1s 24ms/step - loss: 0.2423 - accur
acy: 0.8863 - val loss: 0.3625 - val accuracy: 0.8525
Epoch 49/100
23/23 [============ ] - 0s 20ms/step - loss: 0.2432 - accur
acy: 0.8869 - val loss: 0.3554 - val accuracy: 0.8566
Epoch 50/100
23/23 [============== ] - 0s 20ms/step - loss: 0.2404 - accur
acy: 0.8860 - val loss: 0.3577 - val accuracy: 0.8586
Epoch 51/100
23/23 [============== ] - 1s 24ms/step - loss: 0.2410 - accur
acy: 0.8858 - val loss: 0.3585 - val accuracy: 0.8525
Epoch 52/100
acy: 0.8866 - val loss: 0.3752 - val accuracy: 0.8463
Epoch 53/100
23/23 [============== ] - 1s 23ms/step - loss: 0.2374 - accur
acy: 0.8885 - val loss: 0.3613 - val accuracy: 0.8607
Epoch 54/100
acy: 0.8877 - val loss: 0.3611 - val accuracy: 0.8607
Epoch 55/100
23/23 [============ ] - 0s 20ms/step - loss: 0.2415 - accur
acy: 0.8870 - val loss: 0.3674 - val accuracy: 0.8627
Epoch 56/100
23/23 [============= ] - 1s 23ms/step - loss: 0.2381 - accur
acy: 0.8884 - val loss: 0.3754 - val accuracy: 0.8443
```

```
Epoch 57/100
23/23 [============== ] - 0s 21ms/step - loss: 0.2351 - accur
acy: 0.8898 - val loss: 0.3750 - val accuracy: 0.8545
Epoch 58/100
acy: 0.8886 - val loss: 0.3584 - val accuracy: 0.8689
Epoch 59/100
23/23 [============= ] - 1s 23ms/step - loss: 0.2338 - accur
acy: 0.8871 - val loss: 0.3633 - val accuracy: 0.8545
Epoch 60/100
23/23 [============== ] - Os 20ms/step - loss: 0.2367 - accur
acy: 0.8881 - val loss: 0.3700 - val accuracy: 0.8607
Epoch 61/100
23/23 [============== ] - 0s 20ms/step - loss: 0.2372 - accur
acy: 0.8886 - val loss: 0.3648 - val accuracy: 0.8586
Epoch 62/100
acy: 0.8870 - val loss: 0.3596 - val accuracy: 0.8422
Epoch 63/100
23/23 [============== ] - Os 21ms/step - loss: 0.2378 - accur
acy: 0.8886 - val loss: 0.3546 - val accuracy: 0.8525
Epoch 64/100
23/23 [============= ] - 0s 21ms/step - loss: 0.2340 - accur
acy: 0.8902 - val loss: 0.3504 - val accuracy: 0.8648
Epoch 65/100
23/23 [============== ] - 1s 23ms/step - loss: 0.2316 - accur
acy: 0.8906 - val loss: 0.3530 - val accuracy: 0.8586
Epoch 66/100
23/23 [============== ] - 0s 21ms/step - loss: 0.2343 - accur
acy: 0.8877 - val_loss: 0.3518 - val_accuracy: 0.8525
Epoch 67/100
23/23 [============== ] - 0s 20ms/step - loss: 0.2368 - accur
acy: 0.8883 - val loss: 0.3570 - val accuracy: 0.8607
Epoch 68/100
acy: 0.8893 - val loss: 0.3680 - val accuracy: 0.8525
23/23 [============== ] - 0s 21ms/step - loss: 0.2325 - accur
acy: 0.8887 - val loss: 0.3487 - val accuracy: 0.8586
Epoch 70/100
23/23 [============== ] - 0s 21ms/step - loss: 0.2320 - accur
acy: 0.8925 - val_loss: 0.3654 - val_accuracy: 0.8566
Epoch 71/100
acy: 0.8922 - val loss: 0.3515 - val accuracy: 0.8586
Epoch 72/100
23/23 [============ ] - 0s 19ms/step - loss: 0.2291 - accur
acy: 0.8949 - val loss: 0.3566 - val accuracy: 0.8586
Epoch 73/100
23/23 [============== ] - 1s 23ms/step - loss: 0.2373 - accur
acy: 0.8895 - val loss: 0.3567 - val accuracy: 0.8750
Epoch 74/100
23/23 [============== ] - 1s 23ms/step - loss: 0.2353 - accur
acy: 0.8894 - val loss: 0.3384 - val accuracy: 0.8750
```

```
acy: 0.8891 - val loss: 0.3475 - val accuracy: 0.8730
Epoch 76/100
23/23 [============ ] - 1s 22ms/step - loss: 0.2380 - accur
acy: 0.8877 - val loss: 0.3623 - val accuracy: 0.8586
Epoch 77/100
acy: 0.8900 - val loss: 0.3536 - val accuracy: 0.8750
Epoch 78/100
23/23 [============= ] - 0s 20ms/step - loss: 0.2297 - accur
acy: 0.8911 - val loss: 0.3525 - val accuracy: 0.8607
Epoch 79/100
23/23 [============== ] - 0s 22ms/step - loss: 0.2295 - accur
acy: 0.8912 - val loss: 0.3450 - val accuracy: 0.8750
23/23 [============ ] - 1s 23ms/step - loss: 0.2270 - accur
acy: 0.8940 - val loss: 0.3513 - val accuracy: 0.8586
Epoch 81/100
23/23 [============== ] - 1s 22ms/step - loss: 0.2282 - accur
acy: 0.8913 - val loss: 0.3474 - val accuracy: 0.8607
Epoch 82/100
acy: 0.8913 - val loss: 0.3598 - val accuracy: 0.8627
Epoch 83/100
23/23 [============== ] - 0s 21ms/step - loss: 0.2333 - accur
acy: 0.8905 - val loss: 0.3557 - val accuracy: 0.8709
Epoch 84/100
acy: 0.8897 - val loss: 0.3537 - val accuracy: 0.8709
Epoch 85/100
acy: 0.8915 - val loss: 0.3460 - val accuracy: 0.8627
23/23 [============= ] - Os 22ms/step - loss: 0.2305 - accur
acy: 0.8921 - val loss: 0.3519 - val accuracy: 0.8586
Epoch 87/100
23/23 [============== ] - 0s 20ms/step - loss: 0.2307 - accur
acy: 0.8926 - val loss: 0.3529 - val accuracy: 0.8709
Epoch 88/100
23/23 [============== ] - 1s 22ms/step - loss: 0.2294 - accur
acy: 0.8934 - val loss: 0.3530 - val accuracy: 0.8689
Epoch 89/100
acy: 0.8885 - val loss: 0.3545 - val accuracy: 0.8668
Epoch 90/100
23/23 [============= ] - 0s 20ms/step - loss: 0.2316 - accur
acy: 0.8918 - val loss: 0.3472 - val accuracy: 0.8627
Epoch 91/100
acy: 0.8904 - val_loss: 0.3605 - val_accuracy: 0.8545
Epoch 92/100
23/23 [============== ] - 0s 21ms/step - loss: 0.2318 - accur
acy: 0.8898 - val loss: 0.3721 - val accuracy: 0.8627
Epoch 93/100
acy: 0.8921 - val loss: 0.3446 - val accuracy: 0.8709
Epoch 94/100
```

```
acy: 0.8919 - val loss: 0.3574 - val accuracy: 0.8648
Epoch 95/100
acy: 0.8935 - val_loss: 0.3698 - val_accuracy: 0.8730
Epoch 96/100
acy: 0.8885 - val loss: 0.3414 - val accuracy: 0.8709
Epoch 97/100
23/23 [============== ] - 0s 22ms/step - loss: 0.2337 - accur
acy: 0.8927 - val loss: 0.3423 - val accuracy: 0.8627
Epoch 98/100
acy: 0.8927 - val loss: 0.3421 - val accuracy: 0.8668
Epoch 99/100
23/23 [============== ] - 0s 20ms/step - loss: 0.2326 - accur
acy: 0.8895 - val loss: 0.3637 - val accuracy: 0.8525
Epoch 100/100
acy: 0.8945 - val_loss: 0.3658 - val_accuracy: 0.8566
```

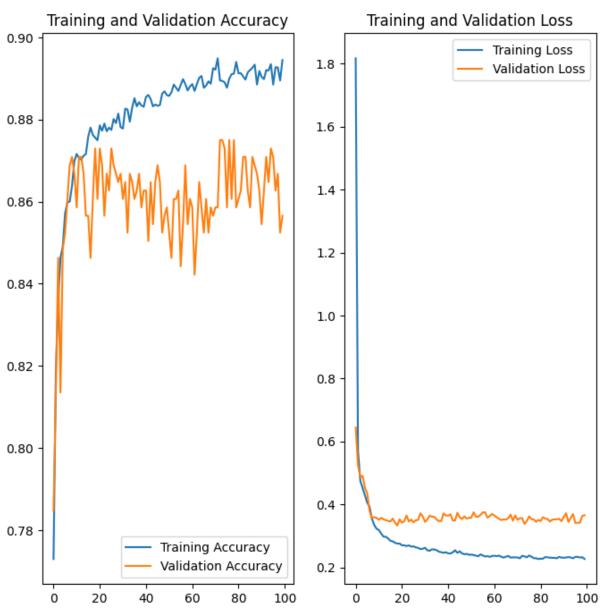
To avoid training and retraining the model, I'll now save this model with these results.

Visualizing the Training Results

Below we can see a visualization of how loss and accuracy evolved over the course of training. Surprisingly, performance is substantially better in this model!

```
acc = nep eng history.history['accuracy']
In [ ]:
        val acc = nep eng history.history['val accuracy']
        loss = nep eng history.history['loss']
        val loss = nep eng history.history['val loss']
        epochs range = range(100)
        plt.figure(figsize=(8, 8))
        plt.subplot(1, 2, 1)
        plt.plot(epochs range, acc, label='Training Accuracy')
        plt.plot(epochs range, val acc, label='Validation Accuracy')
        plt.legend(loc='lower right')
        plt.title('Training and Validation Accuracy')
        plt.subplot(1, 2, 2)
        plt.plot(epochs range, loss, label='Training Loss')
        plt.plot(epochs range, val loss, label='Validation Loss')
        plt.legend(loc='upper right')
```

```
plt.title('Training and Validation Loss')
plt.show()
```



English Language Text Generator

For the text generator, I will create a mini-GPT model for text generation uses the KerasNLP library. This model will be trained on the simplebooks-92 dataset. This dataset uses a simplified English vocabulary. This is useful both for training purposes and for creating generated output that is readily understandable by individuals who do not speak English as a first language.

To create this model, I drew on the Keras suggested tutorial for text generation that can be found here:

https://keras.io/examples/generative/text_generation_gpt/

Settings

Below I've selected some key hyperparameters. Particularly notable here is the minimum traing sequence length. This sets the smallest number of tokens that will be examined by the model during training. We want this number to be large enough that the model is not attempting to train on single words or brief phrases which may eat up training time while providing little in the way of performance improvements.

Additionally, NUM_TOKENS_TO_GENERATE decides how many tokens will appear in the output of the model. I've set this value to be relatively low. There are two reasons for this.

First, in testing, shorter outputs were more likely to be coherent sentences that expressed a single proposition. Given that outputs are meant to be translated and understood, coherence is extremely important.

Second, because the output sentences are meant to be translated and learned by a student, we want them to be short and simple enough to be readily understood and learned, even by a beginner.

```
In []: BATCH_SIZE = 64
SEQ_LEN = 128
MIN_TRAINING_SEQ_LEN = 450

EMBED_DIM = 256
FEED_FORWARD_DIM = 256
NUM_HEADS = 3
NUM_LAYERS = 2
VOCAB_SIZE = 5000

EPOCHS = 50

NUM_CHAR_TO_GENERATE = 50
```

Load Simplebooks Data

Tokenizer for Generator

Here I've defined the vocabulary for the model. This vocabulary is made up of words ('tokens') from the dataset that the model needs to be able to represent and understand.

PAD, UNK, BOS represent padding, unknown, and beginning-of-sentence. These tokens are set aside as non-words for our purposes.

I've then loaded in KerasNLP's tokenizer and used it to preprocess the data for training. This strips out punctuation, sets every word to be all lowercase and then assigns a unique integer to each word. This allows the model to train on the data as quantified data.

Although I've already created a vocabulary and tokenizer for the translator models above, because the set of English sentences in the translation data set is so small, it is necessary to create new ones for the generator model.

Note, if you are re-running this notebook, the vocabulary computation is quite slow.

```
In [ ]: vocab = keras nlp.tokenizers.compute word piece vocabulary(
            raw train ds,
            vocabulary_size = VOCAB_SIZE,
            lowercase = True,
            reserved tokens = ["[PAD]", "[UNK]", "[BOS]"],
In [ ]: gen tokenizer = keras nlp.tokenizers.WordPieceTokenizer(
            vocabulary=vocab,
            sequence length=SEQ LEN,
            lowercase=True,
In [ ]: start packer = keras nlp.layers.StartEndPacker(
            sequence length=SEQ LEN,
            start value=gen tokenizer.token to id("[BOS]")
        def preprocess(inputs):
            outputs = gen tokenizer(inputs)
            features = start packer(outputs)
            labels = outputs
            return features, labels
        train ds = raw train ds.map(tf.autograph.experimental.do not convert(preproc
        val ds = raw val ds.map(tf.autograph.experimental.do not convert(preprocess)
```

Constructing the Model

Now we can actually create the model. This model is another Auto-encoder.

This model is very similar architecturally to the translator model. Tokens are passed to an input layer and then to an embedding layer which accounts for both the number

associated with the token as well as the position of the token in its sentential context. This then downsamples the sentence to an embedding.

These embeddings are then passed to a decoder layer that attempts to generate a sentence by upsampling from the embedding.

Below, I've compiled the model. As before, the loss function is Sparse Categorical Crossentropy. Again, natural language is almost always going to be sparse data so this is the natural choice.

Here I've used Adam as the optimizer. The smallness of the dataset isn't really a concern here so I've chosen Adam for it's ability to produce higher accuracy, even though the computation costs are a little higher.

```
In [ ]: generator.compile(optimizer='adam', loss=loss_fn, metrics=[])
```

Model Summary

```
In [ ]: generator.summary()
```

33 of 41 2023-08-18, 2:32 p.m.

Model: "model 3"

Layer (type)	Output Shape	Param #
input_2 (InputLayer)	[(None, None)]	0
<pre>token_and_position_embeddi ng_3 (TokenAndPositionEmbe dding)</pre>	(None, None, 256)	1312768
<pre>transformer_decoder_3 (Tra nsformerDecoder)</pre>	(None, None, 256)	394749
<pre>transformer_decoder_4 (Tra nsformerDecoder)</pre>	(None, None, 256)	394749
dense_16 (Dense)	(None, None, 5000)	1285000

Total params: 3387266 (12.92 MB)
Trainable params: 3387266 (12.92 MB)
Non-trainable params: 0 (0.00 Byte)

Model Training

Below, I've trained the model. Note that loss is dropping very slowly, which is to say, for a dataset this size the model converges very slowly. For a final product, far more epochs would be needed to achieve satisfactory results.

```
Epoch 2/50
val loss: 3.7982
Epoch 3/50
val loss: 3.7830
Epoch 4/50
val loss: 3.8098
Epoch 5/50
val loss: 3.7441
Epoch 6/50
val_loss: 3.7557
Epoch 7/50
val loss: 3.7628
Epoch 8/50
val loss: 3.7362
Epoch 9/50
val loss: 3.7375
Epoch 10/50
val loss: 3.7417
Epoch 11/50
val loss: 3.7191
Epoch 12/50
val_loss: 3.7226
Epoch 13/50
val loss: 3.6952
Epoch 14/50
val loss: 3.6960
Epoch 15/50
val_loss: 3.7120
Epoch 16/50
val loss: 3.6950
Epoch 17/50
val_loss: 3.7143
Epoch 18/50
val loss: 3.7224
Epoch 19/50
val loss: 3.7073
Epoch 20/50
```

```
val loss: 3.7016
Epoch 21/50
val loss: 3.6911
Epoch 22/50
val loss: 3.7070
Epoch 23/50
val loss: 3.6862
Epoch 24/50
val loss: 3.6918
Epoch 25/50
val loss: 3.6840
Epoch 26/50
val loss: 3.6916
Epoch 27/50
val loss: 3.6916
Epoch 28/50
val loss: 3.7013
Epoch 29/50
val loss: 3.6827
Epoch 30/50
val loss: 3.6793
Epoch 31/50
val loss: 3.6709
Epoch 32/50
val loss: 3.6897
Epoch 33/50
val loss: 3.7001
Epoch 34/50
val loss: 3.6823
Epoch 35/50
val loss: 3.6880
Epoch 36/50
val loss: 3.6953
Epoch 37/50
val loss: 3.6737
Epoch 38/50
val loss: 3.6803
Epoch 39/50
```

```
val loss: 3.6807
  Epoch 40/50
  val loss: 3.6953
  Epoch 41/50
  val_loss: 3.6894
  Epoch 42/50
  val loss: 3.6887
  Epoch 43/50
  val loss: 3.6860
  Epoch 44/50
  val loss: 3.6890
  Epoch 45/50
  val loss: 3.6725
  Epoch 46/50
  val loss: 3.6919
  Epoch 47/50
  val loss: 3.6709
  Epoch 48/50
  val loss: 3.6882
  Epoch 49/50
  val loss: 3.6811
  Epoch 50/50
  val loss: 3.6706
Out[]: <keras.src.callbacks.History at 0x7f908d113a00>
   Like before, I'll save and reload the model for future testing.
In [ ]: # generator.save('models/text-generator.keras')
```

Testing the Text Generator

Let's see what sort of results we're getting from the generator.

In []: generator = tf.keras.models.load model('models/text-generator.keras')

Like for the translator models, we need to select a sampler. The sampler is the algorithm used to pick the next work in our output sequence based on the probabilities that were calculated by the model.

There are a a number of samplers to choose from. I've found the best results, the most

coherent sentences, from using the Top P Sampler.

The Top P Sampler takes a probability, say .9, and selects the most likely sequence of words that sum to that probability. This is different from other samplers which only select the most probable next single word. By using Top P we can get outputs that more closely resemble coherent phrases rather than strings of words that only make sense in context of the couple of words surrounding them.

Below, I've defined a "generate" function which takes in a prompt and returns a sentence generated by our model. Note that the generated outputs need to be cleaned before they are presentable. Like with the translations they are stripped of tokens that simply convey the beginnings and endings of sentences as well as unnecessary whitespace and potentially confusing punctuation.

```
In [ ]: def generate():
            def next(prompt, cache, index):
                logits = generator(prompt)[:, index - 1, :]
                hidden states = None
                return logits, hidden states, cache
            def clean(txt):
                clean txt = txt.numpy()[0].decode("utf-8")
                clean txt = (
                clean txt.replace("[PAD]", "")
                .replace("[START]", "")
                .replace("[END]", "")
                .replace("[BOS]", "")
                 .replace(",", "")
                 .replace("!", ".")
                .replace("?", ".")
                 .replace("\'", "")
                replace("\"", "")
                 .replace(";", "")
                 .replace(":", "")
                clean txt = clean txt[:NUM CHAR TO GENERATE].split(".")[0].strip()
                return clean txt
            prompt = start_packer(gen_tokenizer(['']))
            sampler = keras_nlp.samplers.TopPSampler(p=0.9)
            output tokens = sampler(
                next=next,
                prompt = prompt,
                index=1,
            txt = clean(gen tokenizer.detokenize(output tokens))
            return txt
```

Ok, now let's pass a prompt to the model and see what we get!

```
In [ ]: txt = generate()
    print(f"Generated text: \n{txt}\n")
```

Generated text: the bulk of a wagon had departed in front of his

Conclusion

We now have functions that translate from English to Nepali and from Nepali to English. We also have the ability to generate novel sentences in English. Let's take a look again at some examples.

Example Translations

```
** Example 0 **
no one brought us anything.
तिमिसँग कुनुैख्ने छ ।

** Example 1 **
i didn't want to go, but i did.
मलाई मौले सक्षध्नेष्ठ ।

** Example 2 **
sharkskin can be used like the skin of other animals.
सुनकको लागि अपष्ट वाद्डा हो ।

** Example 3 **
tom just wanted to be helpful.
टम जसरी चाहन्छु ।

** Example 4 **
this tea is too bitter.
यो अक्बि धेरै राउँदै छ ।
```

Text Generation Examples

Translating Generated Text

Now let's look at translations of generated text.

```
In []: for i in range(5):
    input_sentence = generate()
    translated = eng_nep_translate(tf.constant([input_sentence]))
    translated = translated.numpy()[0].decode("utf-8")
    translated = (
        translated.replace("[PAD]", "")
        .replace("[START]", "")
```

```
.replace("[END]", "")
          .strip()
     )
     print(f"** Translation of Generated Text {i} **")
     print(input sentence)
     print(translated)
     print()
** Translation of Generated Text 0 **
as time went on their stairs they made up their
जर्म्रो विदाको राजधानी रोद्याम्दिन मन पर्छ ।
** Translation of Generated Text 1 **
i wish it had to turn to the assault of trodden
मैले के यो शंकाउनुइनुभगी ल्याप्त पीचको उठाबीमा हुनुचा घाइते भए ।
** Translation of Generated Text 2 **
several men were very anxious to know all the fac
मलाई स्पष्ट रूपष्ट प्राकुक आगे केही फावना छैन ।
** Translation of Generated Text 3 **
his majesty spoke of his angry words and was rea
आफ्नो गर्मीमा धेरै रोधें रात्तिताई दरस्ताता त मारिए वा घाइते भए ।
```

So we've done it! This is a great starting place for an AI language tutor. We can translate a into and out of a new language and generate novel sentences for translation and study.

** Translation of Generated Text 4 **

ब्रसेल पुचाउको नजिम्क पवित्रा चाहन्छु ।

by daybreak sir john began to webbergs on the mo

Obviously these models are only a starting point. For a more robust product we will need much more robust data and significantly more training time. This is unfortunately a problem that plagues work on global minority languages. But, hopefully with tools like these we can start to move in the right direction!

41 of 41